



# Schedules and Serializability

Database System Concepts, 6<sup>th</sup> Ed.

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Chapter 14: Transactions

- Concurrent Executions
- Serializability
- Implementation of Isolation



# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system.  
Advantages are:
  - **increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **reduced average response time** for transactions: short transactions need not wait behind long ones.
- **Concurrency control schemes** – mechanisms to achieve isolation
  - that is, to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database
    - ▶ Will study in Chapter 16, after studying notion of correctness of concurrent executions.



# Schedules

- **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - a schedule for a set of transactions must consist of all instructions of those transactions
  - must preserve the order in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a commit instructions as the last statement
  - by default transaction assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- A serial schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



# Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is **not a serial schedule**, but it is **equivalent** to Schedule 1.

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )
read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

In Schedules 1, 2 and 3, the sum **A + B is preserved**.



# Schedule 4

- The following concurrent schedule **does not preserve** the value of  $(A + B)$ .

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )
write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	$B := B + temp$ write ( $B$ ) commit



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency.
- Thus serial execution of a set of transactions preserves database consistency.
- A (possibly concurrent) **schedule is serializable if it is equivalent to a serial schedule**. Different forms of schedule equivalence give rise to the notions of:
  1. **conflict serializability**
  2. **view serializability**



# *Simplified view of transactions*

- We **ignore** operations **other** than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our **simplified schedules** consist of only **read** and **write** instructions.



# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .
  1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  **don't conflict**.
  2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They **conflict**.
  3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They **conflict**
  4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They **conflict**
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.



# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule



# Conflict Serializability (Cont.)

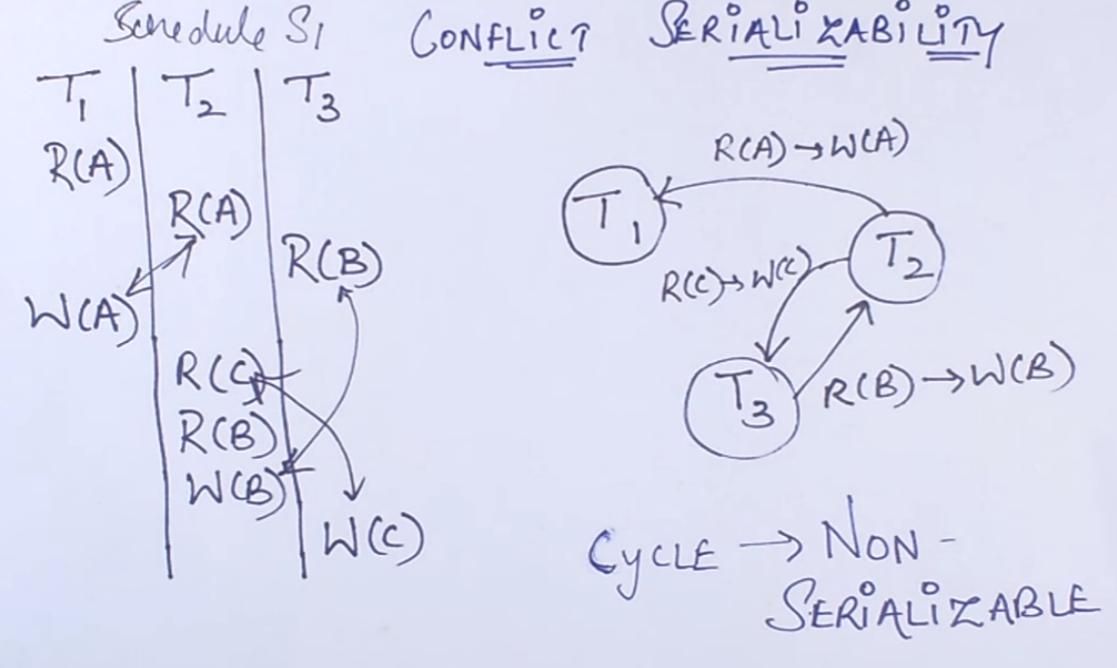
- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions. Therefore Schedule 3 is conflict serializable.

$T_1$	$T_2$
read ( $A$ ) write ( $A$ )	read ( $A$ ) write ( $A$ )
read ( $B$ ) write ( $B$ )	read ( $B$ ) write ( $B$ )

Schedule 3

$T_1$	$T_2$
read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )	read ( $A$ ) write ( $A$ ) read ( $B$ ) write ( $B$ )

Schedule 6



## Serializability: A Broader View

Serializability



Conflict  
Serializable?

Yes

Serializable

Yes

No

View  
Serializable?

No

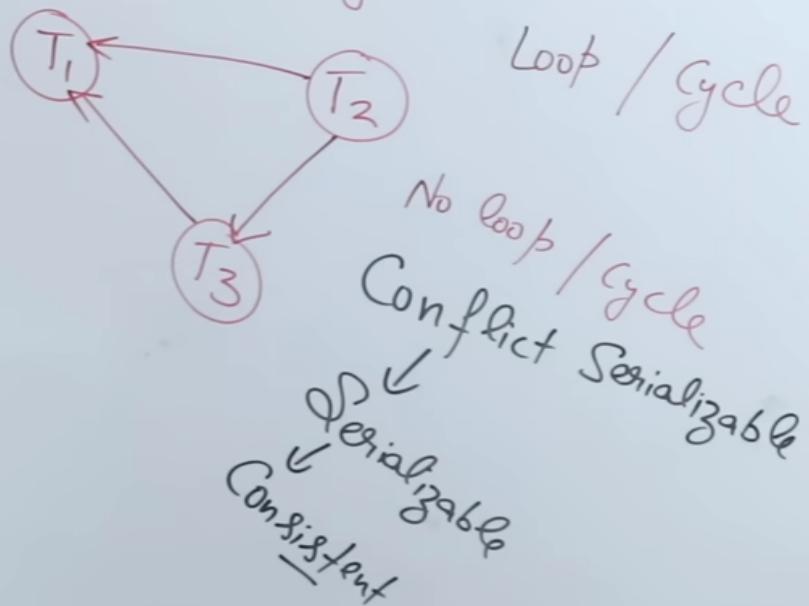
Non-  
Serializable

## Conflict S Serializability

$R-W$   
 $W-R$   
 $W-W$

	$T_1$	$T_2$	$T_3$
	$R(x)$		
$w(y)$		$R(y)$ $R(x)$	
	$R(y)$ $R(z)$		
		$w(y) <^R_R y$	
$w(z)$			$w(y) <^R_R y$
	$R(z)$ $w(x)$ $w(z)$		

- Check Conflict pairs in other transactions and draw edges
- Precedence graph



$(T_2) \rightarrow T_1 \rightarrow (T_3)$

Serial Schedule	$T_1$	$T_2$	$T_3$
		$r(B)$	
		$r(A)$	
	$w(B)$		
	$w(A)$		
			$w(A)$

View Serializability & View Serializable Schedule

IR: by  $T_2$  on data item  $\overset{A}{(A)}$   
 $\overset{B}{(B)}$

FW: by  $T_1$  on data item  $\overset{B}{(B)}$   
by  $T_3$  on data item  $\overset{A}{(A)}$

WR: Not present so need to check

### Example :

Consider the following Non-Serial Schedule  $S_1$ :

$T_1$	$T_2$	$T_3$
	$r(B)$	
	$r(A)$	
$w(B)$		
$w(A)$		
	$w(A)$	
		$w(A)$

→ Not Conflict  
Serializable

1. Initial Read
2. Final Write
3. Write-Read

is Schedule  $S_1$  View-Serializable  
Schedule?



# Anomalies with Interleaved Execution

- Reading Uncommitted Data (WR Conflicts, “**dirty reads**”):

$T_1$	$T_2$
read ( $A$ )	
write ( $A$ )	read ( $A$ )
	write ( $A$ )
	C
read ( $B$ )	
write ( $B$ )	
Abort	

Occurs when a transaction reads data that has been modified by another transaction, but not yet committed



# Anomalies with Interleaved Execution

## □ Unrepeatable Reads (RW Conflicts):

$T_1$	$T_2$
read ( $A$ )	
	read ( $A$ ) write ( $A$ ) $C$
read ( $A$ ) write ( $B$ ) $C$	

$T_1$	$T_2$
$R(A)$	
$W(A)$	
	$R(A)$
	$W(A)$
	$R(B)$
	$W(B)$
	<i>Com.</i>
	$R(B)$
	$W(B)$
	<i>Com.</i>



T2 is reading data which is written by T1 before committing.

WR conflict | CSEstack.org

A user reads two different values for the same record because another user updated the record in between the two reads.



# Anomalies (Continued)

- Overwriting Uncommitted Data (WW Conflicts):

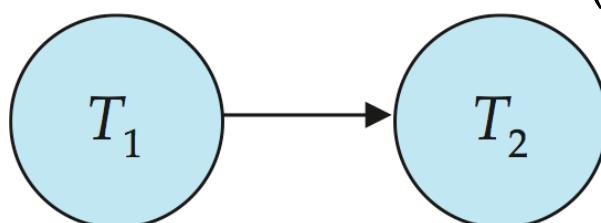
$T_1$	$T_2$
write ( $A$ )	write ( $A$ ) write ( $B$ ) $C$
write ( $B$ ) $C$	

also known as overwriting uncommitted data

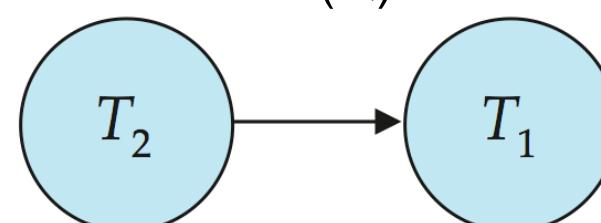


# Precedence graph for schedule 1 and schedule 2.

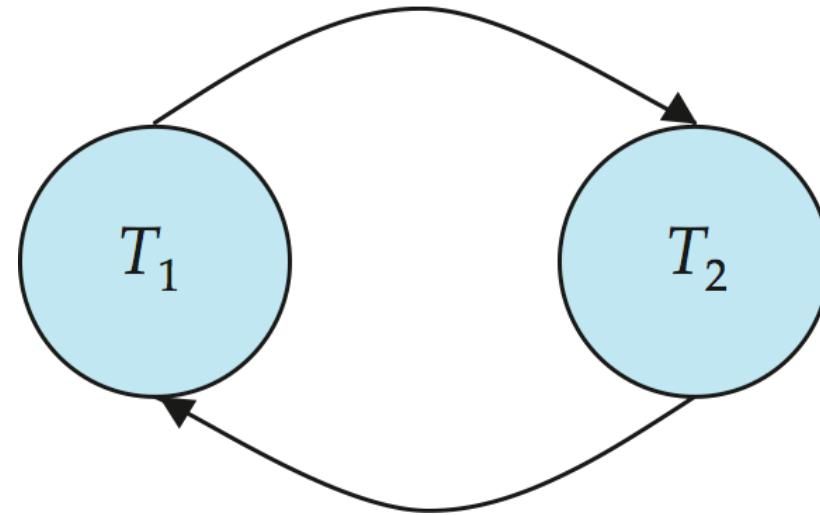
- Simple and efficient method for determining conflict serializability of a schedule.
- Consider a schedule S. We construct a directed graph, called a **precedence graph**, from S.
- This graph consists of a pair  $G = (V, E)$ , where  $V$  is a set of vertices and  $E$  is a set of edges.
- The set of edges consists of all edges  $T_i \rightarrow T_j$  for which one of three conditions holds:
  1.  $T_i$  executes write(Q) before  $T_j$  executes read(Q).
  2.  $T_i$  executes read(Q) before  $T_j$  executes write(Q).
  3.  $T_i$  executes write(Q) before  $T_j$  executes write(Q).

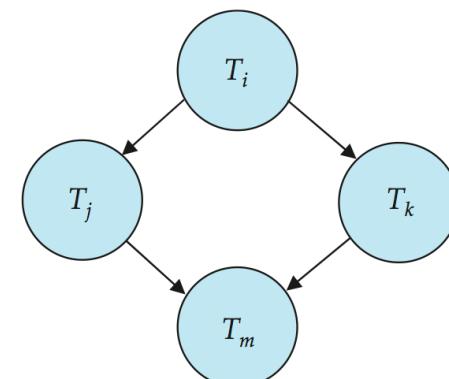


(a)

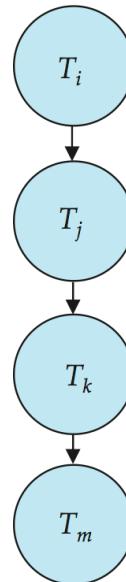


(b)

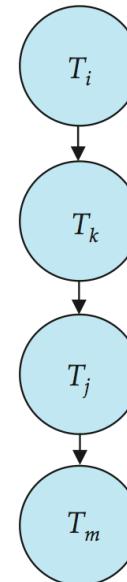




(a)



(b)



(c)



# Figure 14.13

$T_1$	$T_5$
read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $B$ ) $B := B - 10$ write ( $B$ )
read ( $B$ ) $B := B + 50$ write ( $B$ )	
	read ( $A$ ) $A := A + 10$ write ( $A$ )



# View Serializability

- Consider two schedules  $S$  and  $S'$ , where the same set of transactions participates in both schedules. The schedules  $S$  and  $S'$  are said to be **view equivalent** if three conditions are met:
  - For each data item  $Q$ , if transaction  $T_i$  reads the initial value of  $Q$  in schedule  $S$ , then transaction  $T_i$  must, in schedule  $S'$ , also read the initial value of  $Q$ .
  - For each data item  $Q$ , if transaction  $T_i$  executes  $\text{read}(Q)$  in schedule  $S$ , and if that value was produced by a  $\text{write}(Q)$  operation executed by transaction  $T_j$ , then the  $\text{read}(Q)$  operation of transaction  $T_i$  must, in schedule  $S'$ , also read the value of  $Q$  that was produced by the same  $\text{write}(Q)$  operation of transaction  $T_j$ .
  - For each data item  $Q$ , the transaction (if any) that performs the final  $\text{write}(Q)$  operation in schedule  $S$  must perform the final  $\text{write}(Q)$  operation in schedule  $S'$ .



# View Serializability (Cont.)

- Schedule 5 is view equivalent to the serial schedule  $\langle T_{27}, T_{28}, T_{29} \rangle$ , since the one  $\text{read}(Q)$  instruction reads the initial value of  $Q$  in both schedules and  $T_{29}$  performs the final write of  $Q$  in both schedules.
- Every conflict-serializable schedule is also view serializable, but there are view-serializable schedules that are not conflict serializable.
- In schedule 5, transactions  $T_{28}$  and  $T_{29}$  perform  $\text{write}(Q)$  operations without having performed a  $\text{read}(Q)$  operation.
- Writes of this sort are called **blind writes**.
- Blind writes appear in any view-serializable schedule that is not conflict serializable.

Schedule 4

$T_{27}$	$T_{28}$
$\text{read}(Q)$	
	$\text{write}(Q)$

Schedule 5

$T_{27}$	$T_{28}$	$T_{29}$
$\text{read } (Q)$		
$\text{write } (Q)$	$\text{write } (Q)$	$\text{write } (Q)$



# Transaction Isolation and Atomicity

- Lets see the effect of transaction failures during concurrent execution.
- If a transaction  $T_i$  fails, for whatever reason, we need to undo the effect of this transaction to ensure the atomicity property of the transaction.
- In a system that allows concurrent execution, the atomicity property requires that any transaction  $T_j$  that is dependent on  $T_i$  is also aborted.
- To achieve this, we need to place restrictions on the type of schedules permitted in the system.
  - Non-recoverable Schedules
  - Recoverable Schedules
  - Cascading Rollback
  - Cascadeless Schedules



# Nonrecoverable Schedules

- Consider the partial schedule in which T7 is a transaction that performs only one instruction: read(A).
- We call this a partial schedule because we have not included a commit or abort operation for T6.
- Notice that T7 commits immediately after executing the read(A) instruction. Thus, T7 commits while T6 is still in the active state.
- Now suppose that T6 fails before it commits. T7 has read the value of data item A written by T6 . Therefore, we say that T7 is dependent on T6.
- Because of this, we must abort T7 to ensure atomicity. However, T7 has already committed and cannot be aborted. Thus, we have a situation where it is impossible to recover correctly from the failure of T6.
- This Schedule is an example of a **nonrecoverable schedule**.

$T_8$	$T_9$
<code>read (A)</code> <code>write (A)</code>	
<code>read (B)</code>	<code>read (A)</code> <code>commit</code>



# Recoverable Schedules

- A recoverable schedule is one where, for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- For the example of the previous schedule is said to be recoverable,  $T_7$  would have to delay committing until after  $T_6$  commits.



# Cascading Rollback

- A situation in which the abort of one transaction forces the abort of another transaction to prevent the second transaction from reading invalid data. Also called "cascading rollback".
- If in a schedule, failure of one transaction causes several other dependent transactions to rollback or abort, then such a schedule is called as a **Cascading Schedule or Cascading Rollback or Cascading Abort**.

$T_8$	$T_9$	$T_{10}$
read( $A$ ) read( $B$ ) write( $A$ )  abort		
	read( $A$ ) write( $A$ )	read( $A$ )



# Cascadeless Schedules

- It is desirable to restrict the schedules to those where cascading rollbacks cannot occur. Such schedules are called **cascadeless schedules**.
- Formally, a cascadeless schedule is one where, for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- It is easy to verify that every cascadeless schedule is also recoverable.

**THANK YOU**

